

**Clinical Decision-Making and Pattern Recognition in Healthcare:
Chain Reasoning, Agentic Generative AI, and Pattern-Recognition Methods for Treatment,
Payment, and Operations**

Prarthana Shiwakoti

Prepared for Cotiviti

June 27, 2026

Abstract

This brief examines how chain-of-thought reasoning, agentic generative artificial intelligence, and the core pattern-recognition methods of classification, prediction, inference, clustering, and time-series anomaly detection are reshaping clinical decision-making across the three regulated functions of treatment, payment, and operations. Reviewing peer-reviewed and applied research from 2021 through 2026, it traces a clear trajectory: the frontier has moved away from squeezing additional accuracy out of single models and toward systems that reason transparently, collaborate as multiple specialized agents, and can be trusted and audited inside a clinical or payer workflow. The brief maps these capabilities onto Cotiviti's business, identifying where the company's data assets, analyst network, and recent interoperability investments create defensible advantages, and proposes three concrete, generative-AI-centered strategic options: an agentic payment-integrity layer, real-time anomaly detection as a product, and a payer-side clinical-reviewer copilot. The central argument is that the binding constraint in healthcare AI has shifted from raw accuracy to trustworthiness and governance, and that this shift plays directly to Cotiviti's strengths.

The Topic and Why It Matters

Clinical decision-making is the structured process by which clinicians integrate domain knowledge, prior experience, and patient-specific data into a diagnostic or therapeutic decision under uncertainty, and pattern recognition. The ability to find structure, regularity, and meaningful deviation hidden inside medical data (JMIR, 2023). What makes this moment different is that the technology is no longer a single model answering a single question. The field is shifting toward agentic generative AI: systems that can plan a task, break it into steps, call the right tools, and check their own work before handing back an answer (IEEE Access, 2025; IEEE SCORed, 2025). Underneath that sits chain-of-thought reasoning, which forces the model to show its work the way a careful physician talks through a differential; an advance that matters as much for trust and auditability as it does for accuracy (Medicina, 2024). In this brief, these capabilities are organized through five complementary methods and mapped against the three jobs every healthcare enterprise must do well: treatment, payment, and operations (Frontiers in Digital Health, 2025). The framing is deliberate, because Cotiviti's business lives precisely at the intersection of those three.

Five Methods That Carry the Work

Method	The question it answers	Representative evidence & TPO anchor
Classification	Which category does this belong to?	ECG and disease labeling; flagging a claim as normal vs. suspicious
Prediction	What outcome is coming, and when?	ICU mortality (AUC 0.89–0.97), 30-day readmission, demand
Inference	What is the causal effect of acting?	Estimating an individual patient's treatment effect
Clustering	What natural subgroups exist?	Patient phenotyping; grouping related fraud claims
Time-series anomaly detection	Which pattern just broke?	Vital-sign and billing anomalies—touches all of TPO

Relevant Trends: Where the Evidence Is Heading (2021–2026)

Read across five years of research and a clear storyline emerges: the frontier is no longer about squeezing another point of accuracy out of a lone model, but about making systems that reason, collaborate, and can be trusted in a clinic. The most visible change is architectural. A growing body of work documents a move from monolithic models toward multi-agent collaboration, where specialized agents - a diagnostician, a devil's advocate, a validator; debate a

case much like a tumor board would, and in doing so measurably reduce the anchoring and confirmation biases that trip up both humans and machines (IEEE Access, 2025; JMIR, 2024).

Reasoning has become the real differentiator. When models are given room to think step by step, their diagnostic quality jumps: one reasoning model reached 93% accuracy across 100 benchmark cases by working through a structured differential rather than guessing in one shot, and purpose-built reasoning strategies lift performance most on exactly the hard, error-prone questions where help is needed (arXiv, 2025a; IEEE MRAI, 2025). It is a useful reminder that in medicine, how a conclusion is reached is often as important as the conclusion itself.

The second throughline is the fight against hallucination. The dominant answer is to ground generation in verifiable sources, and it works: an open-source agent paired with retrieval-augmented generation produced guideline comparisons with no observed hallucinations across 724 synthetic charts, turning retrieval and knowledge-graph grounding into something close to a design standard (Journal of Clinical Oncology, 2025). On the temporal side, encounter-level transformers now predict intensive-care mortality at areas under the curve of 0.89 to 0.97 while still exposing interpretable features, and compact models can run anomaly detection on a wearable under tight power budgets (JAMIA Open, 2025; IEEE ISDFS, 2026).

Payment integrity, meanwhile, is industrializing. With an estimated 3 to 10% of healthcare spending lost to fraud, waste, and abuse, a 137-study review finds the field still leans on supervised methods but is tilting fast toward deep-learning, graph, and unsupervised approaches—ensemble trees already reach 96% accuracy on claims data (Artificial Intelligence in Medicine, 2024; IEEE GINOTECH, 2025). Clustering is following a similar arc from curiosity to clinical tool: grouping roughly 199,000 atrial-fibrillation patients surfaced five subtypes with sharply different five-year outcomes (European Heart Journal Digital Health,

2025). And for the first time the literature carries real deployment data—an EHR-integrated assistant used in 21.7% of 258,106 episodes with 99.5% positive clinician feedback—which tells us the hard problems are now adoption, localization, and workflow, not just model performance (BMJ Digital Health, 2026).

Opportunities

The upside falls into three buckets, and each one rewards a company that already owns data and analysts rather than one starting from scratch. First, these systems can scale expertise where it is scarce: agentic systems can push specialist-grade assessment into settings short on specialists, and can let non-experts run analyses that once required a machine-learning team, compressing cost and easing compliance along the way (Semantic Scholar, 2025; IEEE BigData, 2025). Second, the economics are concrete, not hypothetical—fewer improper payments, cleaner reimbursement cycles, and operational wins such as a predict-then-optimize scheduler that cut conflicts by 41% and waiting times by 18% in pilots (Scientific Reports, 2026; JAMIA Open, 2025). Third, and most strategically, these methods can make care proactive and transparent at the same time: confidence-aware outputs and visible reasoning build trust, while on-device anomaly detection shifts the model of care from reacting to a crisis toward catching it early (arXiv, 2024; IEEE ICIDCA, 2025).

Threats and Risks

The same power cuts the other way, and an honest brief has to say so. Hallucination remains the headline risk: on a physician-validated benchmark of 4,695 samples, even a leading model produced incorrect content 29.1% of the time, with some open models far worse—numbers that make retrieval grounding, uncertainty estimates, and human oversight non-negotiable rather than optional (AAAI, 2026). Bias and opacity are close behind, because a

model can quietly inherit the skew in its training data and, in payment settings, even mask the very fraud it is meant to catch (AAAI, 2026; IJRSS, 2025). Trust erodes when clinicians cannot see why a system decided what it did, which pushes them toward both over- and under-reliance (BMJ Digital Health, 2026). Regulation has not caught up either—frameworks built for static devices struggle with systems that learn and adapt (SSRN, 2024)—and handing autonomy to multiple cooperating agents introduces new, emergent failure modes that demand layered guardrails (IEEE Access, 2025; arXiv, 2024). The strategic read is straightforward: whoever solves governance, not just accuracy, wins this market.

Cotiviti: Company Fit and Strategic Context

What makes this topic feel tailor-made for Cotiviti is that the company already sits exactly where these methods create value. It is a leader in payment integrity, risk adjustment, quality and Stars (HEDIS), and consumer engagement, with visibility across more than 300 million members and a human-in-the-loop model that pairs machine learning with thousands of clinicians and content experts (Cotiviti, n.d.)—precisely the oversight the research says is mandatory. Company leadership has described healthcare as a coordination problem, and that framing maps neatly onto multi-agent AI, whose entire purpose is coordination. The Caspian Insights platform already holds longitudinal data on roughly 40% of Americans, a data moat most competitors simply cannot replicate (BusinessWire, 2020). The topic is also part of Cotiviti’s own history: its 2020 COVID-19 Outbreak Tracker combined classification, pattern recognition, and time-series anomaly detection over ICD-10 and CPT signals to anticipate outbreaks (Cotiviti, 2020). Recent strategy only sharpens the fit—the 2025 Edifecs acquisition added FHIR and EDI interoperability and an agentic AI platform for value-based payment (Cotiviti, 2025; GeekWire, 2025), a 2026 product vision recasts AI as an operational intelligence

layer that spots pipeline anomalies and recommends fixes (Cotiviti, 2026), the flagship 360 and fraud-waste-abuse Pattern Review already blends AI, expert rules, and credentialed analysts (Cotiviti, n.d.-b), and the platform ranked first in Black Book’s 2026 payer claims survey (Black Book Research, 2026). In short, the building blocks are in place; the question is where to point them next.

Three Options for Cotiviti to Explore

Each of the following options takes the report’s named methods—chain reasoning, agentic generative AI, classification, prediction, inference, clustering, and time-series anomaly detection—and points them at something Cotiviti already owns, so the research base turns into a concrete, fundable thesis rather than a wish list.

Option 1: Agentic Payment-Integrity and Fraud-Waste-Abuse Layer

Cotiviti’s 360 Pattern Review and fraud-waste-abuse practice is its most mature, success-fee business, blending machine-learning models, expert-authored rules, and credentialed human analysts, and it was ranked first in the Black Book 2026 payer claims survey (Cotiviti, n.d.-b). The work today still leans heavily on rule-and-score logic: analysts spend long stretches manually gathering evidence, the flags they receive behave like a black box that is hard to defend, and genuinely novel or collusive schemes slip through the seams (IEEE GINOTECH, 2025). A multi-agent system built on the Caspian data foundation could change that—a retrieval agent assembles the relevant claims and history, a reasoning agent powered by a large language model produces a plain-English rationale for why a case looks suspicious, a validation agent checks that reasoning against policy using retrieval-augmented generation, and a graph model surfaces provider-network collusion that pairwise rules miss (IEEE Access, 2025). An orchestrator then hands the analyst a finished, explained case with a confidence score attached.

The effect is that reviewers spend their time judging completed cases rather than building them, the chain-of-thought rationale turns an opaque flag into something auditable, and the combination of graph signals and generative reasoning catches the sophisticated fraud that today's rules overlook—all while a human keeps the final say (JMIR, 2024).

Option 2: Real-Time Anomaly Detection as a Product

Most of Cotiviti's anomaly work is retrospective or runs in prepay batches, even though the company has proven it can operate in real time before—its 2020 COVID-19 Outbreak Tracker combined classification with time-series anomaly detection over live ICD-10 and CPT streams, and Caspian already touches data for roughly forty percent of Americans (BusinessWire, 2020). The structural weakness is the “pay and chase” pattern: by the time a problem surfaces in a retrospective batch, the money has usually already gone out the door, and that real-time muscle was a one-off response rather than a repeatable offering. Productizing time-series anomaly detection into an always-on streaming service would close that gap—streaming unsupervised detectors flag deviations the moment they appear, a generative explanation layer turns each raw anomaly into a readable alert, and tunable thresholds let operators trade off review volume against false positives (IJCCS, 2021). This shifts the catch from after payment to before it and converts a one-time capability into a durable product, while the generative explanations make the output usable by operations staff who are not data scientists (IEEE ICIDCA, 2025).

Option 3: Generative-AI Clinical-Reviewer Copilot

Cotiviti runs a large human-in-the-loop workforce—thousands of clinicians and content experts who manually read charts, claims, and records to validate coding for risk adjustment and HCC capture, confirm quality measures such as HEDIS, and support payment-integrity reviews (Cotiviti, n.d.). That manual chart review is slow, cognitively heavy, error-prone, and hard to

scale, since a reviewer may scroll through dozens of pages to confirm a single code. A generative-AI copilot, deliberately built for the payer-side reviewer rather than the bedside clinician, could carry most of that load: it auto-structures the chart and associated claims into a SOAP-style summary of the subjective, objective, assessment, and plan; it generates a quick visual summary showing the condition timeline, risk flags, active medication list, and open care gaps; and a recommendation agent then suggests supportable actions—likely HCC and diagnosis codes with the exact chart evidence cited, open HEDIS and Stars gaps, adverse drug-interaction or adherence concerns spotted across claims, and member-outreach next-best-actions (JAMIA Open, 2025). Because the whole copilot is grounded in the chart and the policy library through retrieval-augmented generation, every recommendation links back to its source and the human reviewer always makes the final call (JMIR, 2024). Hours of reading collapse into minutes of reviewing a structured, visual, pre-summarized case; the evidence-cited, human-approved design preserves both accuracy and auditability; the system leans directly on assets Cotiviti already has; and it extends naturally into member-facing care-gap closure (Scientific Reports, 2026).

Bottom line

Across five years of evidence the binding constraint has quietly moved from raw accuracy to trustworthiness, governance, and real deployment. That shift happens to play directly to Cotiviti's strengths: a human-in-the-loop network, the Caspian data asset, and Edifecs interoperability together let the company treat hallucination, bias, and compliance as engineering requirements to be solved—and, in doing so, to lead on agentic, anomaly-aware healthcare AI.

References

- Applied Sciences. (2023). Deep learning for ECG: A systematic review. Applied Sciences, MDPI. <https://www.mdpi.com/2076-3417/13/8/4964>
- Artificial Intelligence in Medicine. (2024). Fraud detection: A systematic review. Artificial Intelligence in Medicine. <https://linkinghub.elsevier.com/retrieve/pii/S09333365724003038>
- arXiv. (2024). LLM risks and AI guardrails. arXiv. <https://arxiv.org/abs/2406.12934>
- arXiv. (2025a). Medical reasoning in DeepSeek R1. arXiv. <https://arxiv.org/html/2504.00016v1>
- Association for the Advancement of Artificial Intelligence. (2026). Medical hallucination benchmark (MHB). AAAI. <https://ojs.aaai.org/index.php/AAAI/article/view/41243>
- Black Book Research. (2026). 2026 payer claims platform rankings. Black Book Research. <https://blackbookmarketresearch.com/>
- BMC Medical Informatics and Decision Making. (2025). Explainable unsupervised anomaly detection. BMC MIDM. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11720628/>
- BMJ Digital Health. (2026). EHR-integrated generative AI clinical decision support in Kenya. BMJ Digital Health. <https://bmjdigitalhealth.bmj.com/lookup/doi/10.1136/bmjdhai-2025-000207>
- BusinessWire. (2020). Cotiviti launches Caspian Insights platform. BusinessWire. <https://www.businesswire.com/news/home/20200615005267/en/>
- Cotiviti. (2020). COVID-19 outbreak tracker. <https://www.cotiviti.com/insights>
- Cotiviti. (2025). Cotiviti completes acquisition of Edifecs. <https://www.cotiviti.com/news>
- Cotiviti. (2026). AI for encounter operations. <https://www.cotiviti.com/insights>
- Cotiviti. (n.d.-b). FWA pattern review. <https://www.cotiviti.com/solutions/payment-integrity>

Cotiviti. (n.d.). Company overview. <https://www.cotiviti.com/about>

Data Science and Engineering. (2023). Sub-extreme anomaly detection for billing review.

Springer DSE. <https://link.springer.com/10.1007/s41019-023-00234-7>

European Heart Journal Digital Health. (2025). Atrial-fibrillation subphenotypes via clustering.

European Heart Journal Digital Health.

<https://academic.oup.com/ehjdh/article/6/4/797/8106708>

Frontiers in Digital Health. (2025). AI with agency in healthcare. Frontiers in Digital Health.

<https://pmc.ncbi.nlm.nih.gov/articles/PMC12092461/>

GeekWire. (2025). Cotiviti to acquire Edifecs. GeekWire. <https://www.geekwire.com/2025/>

Heliyon. (2024). Multi-channel heterogeneous graph fraud detection. Heliyon.

<https://pmc.ncbi.nlm.nih.gov/articles/PMC11061682/>

IEEE Access. (2025). Agentic AI: A survey. IEEE Access.

<https://ieeexplore.ieee.org/document/10849561/>

IEEE BigData. (2025). NimbleLabs: Multi-agent medical data analysis. IEEE International

Conference on Big Data. <https://ieeexplore.ieee.org/document/11401464/>

IEEE GINOTECH. (2025). Intelligent fraud detection. IEEE GINOTECH.

<https://ieeexplore.ieee.org/document/11076870/>

IEEE ICIDCA. (2025). Smart healthcare real-time monitoring. IEEE ICIDCA.

<https://ieeexplore.ieee.org/document/11280423/>

IEEE ISDFS. (2026). CNN-transformer for wearable anomaly detection. IEEE International

Symposium on Digital Forensics and Security.

<https://ieeexplore.ieee.org/document/11458952/>

- IEEE MRAI. (2025). Enhancing medical diagnostic reasoning with chain-of-thought. IEEE MRAI. <https://ieeexplore.ieee.org/document/11135834/>
- IEEE SCORed. (2025). AI agents in healthcare: A survey. IEEE Student Conference on Research and Development. <https://ieeexplore.ieee.org/document/11398875/>
- Indonesian Journal of Computing and Cybernetics Systems. (2021). Anomaly detection in hospital claims. IJCCS. <https://jurnal.ugm.ac.id/ijccs/article/view/68160>
- International Journal of Innovative Research and Scientific Studies. (2025). Feature engineering for fraud detection. IJIRSS. <http://www.ijirss.com/index.php/ijirss/article/view/8858>
- JAMIA Open. (2025). TECO: A transformer encounter-level model for ICU mortality. JAMIA Open. <https://academic.oup.com/jamiaopen/article/doi/10.1093/jamiaopen/ooaf026/8110091>
- Journal of Clinical Oncology. (2025). Open-source agentic AI for breast cancer. JCO. https://ascopubs.org/doi/10.1200/JCO.2025.43.16_suppl.e13685
- Journal of Medical Internet Research. (2023). Toward clinical generative AI. JMIR. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11193080/>
- Journal of Medical Internet Research. (2024). Mitigating cognitive biases via multi-agent LLMs. JMIR. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11615553/>
- Medicina. (2024). Chain-of-thought utilization in nephrology. Medicina. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10819595/>
- Nature Medicine. (2024). Causal machine learning for treatment outcomes. Nature Medicine. <https://arxiv.org/html/2410.08770v1>
- PLOS Digital Health. (2025). Step-by-step causal analysis of EHRs. PLOS Digital Health. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11790099/>

Scientific Reports. (2026). AI-driven hospital human-resource management. Scientific Reports.

<https://www.nature.com/articles/s41598-026-43102-w>

Semantic Scholar. (2025). Agentic AI for neurodegenerative disease. Semantic Scholar.

<https://www.semanticscholar.org/paper/b7359c2a36aaf13ce45dd3806ba7553c35b0bc1a>

SSRN. (2024). AI and the medical device regulation. SSRN.

<https://www.ssrn.com/abstract=4638548>